


Compiling Unikernels into Micro Kernels


Diploma Defense

Lisza Zeidler

Performance vs. Security Trade-off




+ Performance
- No Isolation



- Overhead
+ Strong Isolation


Development and Verification

Unikernel / Shared



- + Simple writing and testing
- + Verifiable

Micro kernel / Isolated



- Hard to develop and test
- Verification hard/impossible

Wanted

Develop

shared memory,
single threaded,
no-isolation,
verified




Deploy


isolated as
needed,
concurrent,
still verified

How to Generalizing the Rewrite ?

```
loop {  
  device.get(...);  
  ip_stack.unpack(...);  
  backend.answer(...);  
  ip_stack.pack(...);  
  device.send(...);  
}
```



Idea - Use a Compiler




Ohua¹:

- sequential → deterministic concurrent
- derives Data Flow Graph
- Backend Integrations provide process + channel implementations

¹Sebastian Ertel, Christof Fetzer, and Pascal Felber. "Ohua: Implicit dataflow programming for concurrent systems". In: *Proceedings of the Principles and Practices of Programming on The Java Platform*. 2015, pp. 51–64.

Problem Solved? → No

```
loop {  
  device.get(...);  
  ip_stack.unpack(...);  
  backend.answer(...);  
  ip_stack.pack(...);  
  device.send(...);  
}
```



States are input and output of methods in data flow graphs


Requirement

State Locality: Isolated components or services should stay in their own runtime isolation

Task: Restructure the program, such that **every state** is used **exactly once**

Status Quo

```
loop {  
  device.get(...);  
  ip_stack.unpack(...);  
  backend.answer(...);  
  ip_stack.pack(...);  
  device.send(...);  
}
```




- program is no valid input
- output program would not meet requirements

Ohua does not support multiple state usage in loops or states in branches

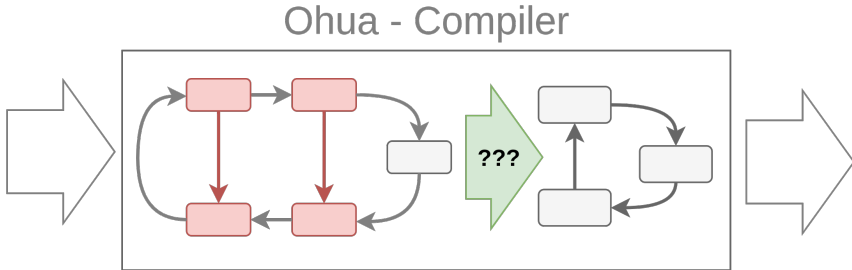
Question 1

```
loop {  
  device.get(...);  
  ip_stack.unpack(...);  
  backend.answer(...);  
  ip_stack.pack(...);  
  device.send(...);  
}
```




How to refactor to a valid input yielding state local programs?

Question 2



Can we teach those refactorings to Ohua?


Concrete Example



- simple server app using `smoltcp`
- M³ OS² as backend

²Nils Asmussen, Michael Roitzsch, and Hermann Härtig. "M3x: Autonomous Accelerators via Context-Enabled Fast-Path Communication". In: *USENIX Annual Technical Conference (ATC)*. Renton, WA, USA: USENIX, July 2019.

Concrete Example



- Goal: Run Device (NIC abstraction), TCP/IP-Stack and Key-Value Store as isolated processes in M³

Concrete Example – App Structure

```
let store = Store::new();  
/* initialization */  
loop {  
    ip_stack.poll(time, &mut device, &mut sockets);  
  
    if let Some(input) = socket.recv(){  
        if socket.can_send() {  
            let outbytes = store.answer(&input);  
            socket.send_slice(&outbytes[..]);  
        }  
    }  
    phy_wait(dev_pointer, ip_stack.delay(time, &sockets));  
}
```

Approach

Refactor a concrete application asking "**What would a compiler do?**"

Insight – Three Kinds of Refactorings

The Good : Solvable, formal Transformations


The Bad: Probably not-solvable, can be included in the Programming Model

The Ugly: Not-solvable, break the Promise

The Good

Refactoring Control Flow – Making States composable


The Good - Restructure Sending Loop




Situation

- outer loop: exchanging messages between store and ip_stack
- inner loop: exchanging packets between ip_stack and device


Plan



start \rightarrow



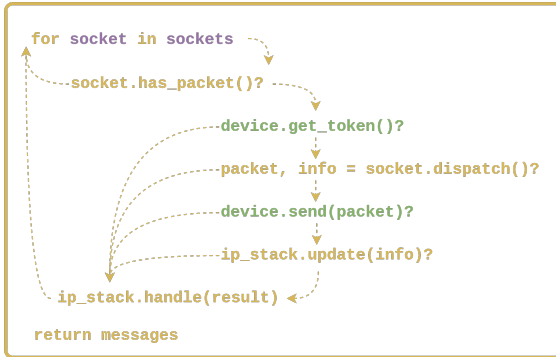
lift device usage into scope →



merge data flow

Inner Loop

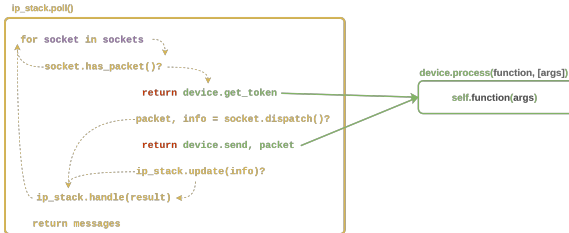
```
ip_stack.poll(&device)
```



We need to call `device` outside of
`ip_stack.poll()`

→ return calls to main scope

Inner Loop



- type of `ip_stack.poll()` and `device.process()` ?
- we can not 'send' functions


Transformation 1 –Defunctionalization³

```
device.get_token  
device.send(packet)  ⇒  
  
enum DeviceCall {  
  GetToken,  
  Send(IPPacket)  
}  
  
device.process(function, [args]) {  
  self.function(args)  
}  ⇒  
  
device.process(call:DeviceCall)-> ???  
{  
  match call {  
    GetToken => self.get_token(),  
    Send(packet) => self.send(packet),  
  }  
}
```

1. define a sum type that represents functions and their arguments
`enum DeviceCall`
2. define a function that *interprets* given values of that type back to function execution

³John C Reynolds. "Definitional interpreters for higher-order programming languages". In: *Proceedings of the ACM annual conference-Volume 2*. 1972, pp. 717–740.

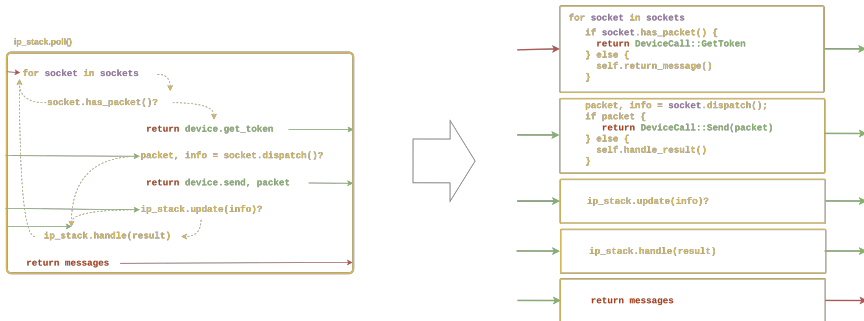
How to send step wise?



We need to return to where we left

- How to call sub-methods?
- How to preserve state of execution?


Transformation 2 – λ -Lifting⁴



turn closures (\approx basic blocks) into top-level functions/ methods

⁴Thomas Johnsson. *Lambda lifting: Transforming programs to recursive equations*. Springer, 1985.

Defunctionalization – Closing the Loop



1. define a sum type that represents lifted functions `enum IPStackCall`
2. define `ip_stack.process` method to interpret the `IPStackCalls`
3. merge data flows

The Bad

Refactoring Reference Usage (in general)

The Bad – Trivial Case

References must not be used as function arguments in scope

Consequence: Programming model assumes for any `function(a, b, c)` in scope `a`, `b`, `c` are passed by value

The Bad – Non-trivial Example

Sending works via tokens:

1. `ip_stack` requests a sending token from `device`
 2. `token` → reference to the device
 3. `token.consume(inner_stack, |buffer| /*closure*/)` is called where `/*closure*/` writes the packet to the buffer provided by the `device`
- ⇒ Highly efficient in shared memory setting but ...

The Bad – Non-trivial Example

Sending works via `tokens`: \Rightarrow Highly efficient in shared memory setting but

- How to split `/*closure*/` to components?
- What to replace the `token` with?
- Are `tokens` useful without references?

Consequence: Refactoring involves dynamic/domain information the compiler does not have

The Ugly

Refactoring System Calls

The Ugly – System/OS calls

```
phy_wait(dev_pointer, ip_stack.delay(...));
```

wait for **File Pointer** until **Timeout**

```
maybe_syscall(a, b, c)
```

- System call?
- Supported by target OS?
- Can we replace by equivalent calls?

The Ugly – System/OS calls

- We can not identify system calls by syntax + static information
- Common practice→ provide annotations^{5, 6, 7}
- What if different code structure is required?

Consequence:

Annotations are an options for simple cases

Complex cases break the Idea

⁵[Hugo Lefeuve](#). “FlexOS: easy specialization of OS safety properties”. In: *Proceedings of the 22nd International Middleware Conference: Doctoral Symposium*. 2021, pp. 29–32.

⁶[Vasily A Sartakov](#), [Lluís Vilanova](#), and [Peter Pietzuch](#). “CubicleOS: a library OS with software componentisation for practical isolation”. In: *Proceedings of the 26th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. 2021, pp. 546–558.

⁷[Carolina Perez Ortega](#). “FlexC: Flexible Compartmentalization Through Automatic Policy Generation”. PhD thesis. Massachusetts Institute of Technology, 2022.

Summary

Question 1 – Can we compile now?

Yes and No:

- add missing syntax support to Ohua
- adapt to M^3 supported types and `device` implementation

Summary

Question 2 – What could Ohua learn?

- + Defunctionalization can be used make states composable
- + λ -Lifting could be used to disentangle states
 - We cannot refactor reference usage
 - We cannot identify system calls

Summary

Question 2.5 – What's the cost?

Rewrite requires:

- owned data types
- copying data for sending
- multiple function calls instead of one in `ip_stack.poll`

⇒ Costs for rewrite on the same OS : TCP Packet throughput (Gb/s) decreased by $\approx 1/3$

Conclusion

Tasks for us:

- Extend Ohuas State Support
- Decide what to compile and what to reject
- Derive Transformations where possible

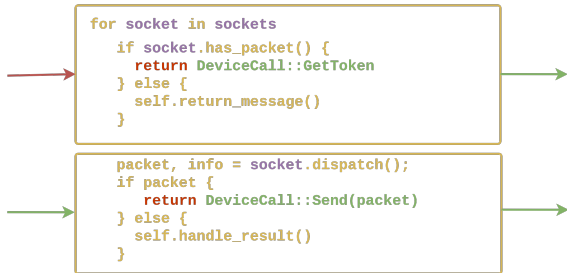
Tasks of the programmer:

- Identify components and make every use explicit
- Stick to the programming model
- Know the target OS's calls and types

And when it's all defined and written ...

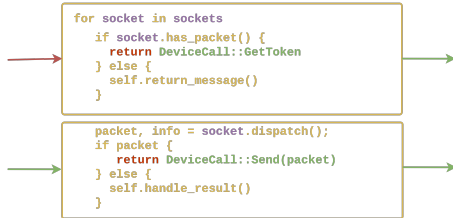


The Bad – λ -Lifting – Problem



- We split a function stack
 - Variables need to get from one Stack to the next
- variables need to be heap allocated or send along with control flow

The Bad – λ -Lifting – Problem



Options

1. Send

- not wanted for state internals
- requires 'sendability'

2. Store in State

- make the socket part of the `ip_stack` state
- heap allocate
- ⇒ No internal cross referencing

The Bad – λ -Lifting – Problem

Splitting function stacks \Rightarrow requirements for reference handling change

Consequence: If this transformation is applied, we need to extend the programming model to 'state internal' code \rightarrow more complex requirements for the programmer

State Threads

```
let b = f(a)
```



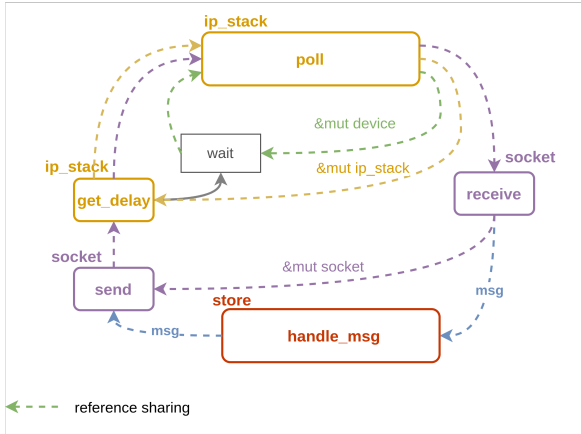
```
let b = obj.handle(a)
```



```
let b = obj.handle(a);  
let c = fun(b);  
let final = obj.process(c);
```



Server Loop: Structure



Problems:

- device not used in scope
- ip_stack used more than once
- communication via shared references
- socket appear as component